

Lending Club

July 8, 2024

```
[1]: #Packages for data reading and manipulation
import pandas as pd
import numpy as np
#for data visualization
import plotly.express as px
import plotly.subplots as sp
import plotly.graph_objs as go
#for encoding
from sklearn.preprocessing import OneHotEncoder
#for data scaling
from sklearn.preprocessing import RobustScaler
#for data splitting
from sklearn.model_selection import train_test_split
#for modelling
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
#for evaluation
from sklearn.metrics import accuracy_score

import warnings
warnings.filterwarnings('ignore')
```

2024-06-03 09:02:36.888380: I tensorflow/core/util/port.cc:110] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.

2024-06-03 09:02:36.926597: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

VOC-NOTICE: GPU memory for this assignment is capped at 1024MiB

2024-06-03 09:02:38.712422: E tensorflow/compiler/xla/stream_executor/cuda/cuda_driver.cc:268] failed call to cuInit: CUDA_ERROR_NO_DEVICE: no CUDA-capable device is detected

```
[2]: #import dataset
df=pd.read_csv("loan_data.csv")
df.head()
```

```
[2]: credit.policy      purpose  int.rate  installment  log.annual.inc  \
0          1  debt_consolidation    0.1189         829.10        11.350407
1          1      credit_card    0.1071         228.22        11.082143
2          1  debt_consolidation    0.1357         366.86        10.373491
3          1  debt_consolidation    0.1008         162.34        11.350407
4          1      credit_card    0.1426         102.92        11.299732

      dti  fico  days.with.cr.line  revol.bal  revol.util  inq.last.6mths  \
0  19.48  737      5639.958333      28854         52.1             0
1  14.29  707      2760.000000      33623         76.7             0
2  11.63  682      4710.000000       3511         25.6             1
3   8.10  712      2699.958333      33667         73.2             1
4  14.97  667      4066.000000       4740         39.5             0

delinq.2yrs  pub.rec  not.fully.paid
0           0         0             0
1           0         0             0
2           0         0             0
3           0         0             0
4           1         0             0
```

```
[3]: #data summary statistics
df.describe()
```

```
[3]: credit.policy      int.rate  installment  log.annual.inc      dti  \
count    9578.000000  9578.000000  9578.000000    9578.000000  9578.000000
mean         0.804970    0.122640    319.089413     10.932117    12.606679
std         0.396245    0.026847    207.071301      0.614813     6.883970
min         0.000000    0.060000    15.670000      7.547502     0.000000
25%         1.000000    0.103900    163.770000     10.558414     7.212500
50%         1.000000    0.122100    268.950000     10.928884    12.665000
75%         1.000000    0.140700    432.762500     11.291293    17.950000
max         1.000000    0.216400    940.140000     14.528354    29.960000

      fico  days.with.cr.line      revol.bal  revol.util  \
count  9578.000000    9578.000000  9.578000e+03  9578.000000
mean   710.846314    4560.767197  1.691396e+04   46.799236
std    37.970537    2496.930377  3.375619e+04   29.014417
min    612.000000    178.958333  0.000000e+00    0.000000
25%    682.000000    2820.000000  3.187000e+03   22.600000
50%    707.000000    4139.958333  8.596000e+03   46.300000
75%    737.000000    5730.000000  1.824950e+04   70.900000
max    827.000000   17639.958330  1.207359e+06  119.000000
```

	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
count	9578.000000	9578.000000	9578.000000	9578.000000
mean	1.577469	0.163708	0.062122	0.160054
std	2.200245	0.546215	0.262126	0.366676
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000
75%	2.000000	0.000000	0.000000	0.000000
max	33.000000	13.000000	5.000000	1.000000

```
[4]: #null values
df.isnull().sum()
```

```
[4]: credit.policy      0
      purpose          0
      int.rate         0
      installment      0
      log.annual.inc   0
      dti              0
      fico             0
      days.with.cr.line 0
      revol.bal        0
      revol.util       0
      inq.last.6mths   0
      delinq.2yrs      0
      pub.rec          0
      not.fully.paid   0
      dtype: int64
```

```
[5]: #data dimensions
df.shape
```

```
[5]: (9578, 14)
```

```
[6]: df.columns
```

```
[6]: Index(['credit.policy', 'purpose', 'int.rate', 'installment', 'log.annual.inc',
          'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
          'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid'],
          dtype='object')
```

```
[7]: #unique values
df[['credit.policy', 'purpose', 'int.rate', 'installment', 'log.annual.inc',
    'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
    'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid']].nunique()
```

```
[7]: credit.policy      2
      purpose          7
      int.rate         249
      installment      4788
      log.annual.inc   1987
      dti              2529
      fico             44
      days.with.cr.line 2687
      revol.bal        7869
      revol.util       1035
      inq.last.6mths   28
      delinq.2yrs      11
      pub.rec          6
      not.fully.paid   2
      dtype: int64
```

```
[8]: # variables data types
      df.dtypes
```

```
[8]: credit.policy      int64
      purpose          object
      int.rate         float64
      installment      float64
      log.annual.inc   float64
      dti              float64
      fico             int64
      days.with.cr.line float64
      revol.bal        int64
      revol.util       float64
      inq.last.6mths   int64
      delinq.2yrs      int64
      pub.rec          int64
      not.fully.paid   int64
      dtype: object
```

```
[9]: #EXPLOTATORY ANALYSIS
```

```
[10]: df.head()
```


```
[10]:   credit.policy  purpose  int.rate  installment  log.annual.inc \
0             1  debt_consolidation    0.1189         829.10    11.350407
1             1    credit_card    0.1071         228.22    11.082143
2             1  debt_consolidation    0.1357         366.86    10.373491
3             1  debt_consolidation    0.1008         162.34    11.350407
4             1    credit_card    0.1426         102.92    11.299732

      dti  fico  days.with.cr.line  revol.bal  revol.util  inq.last.6mths \
```

0	19.48	737	5639.958333	28854	52.1	0
1	14.29	707	2760.000000	33623	76.7	0
2	11.63	682	4710.000000	3511	25.6	1
3	8.10	712	2699.958333	33667	73.2	1
4	14.97	667	4066.000000	4740	39.5	0

	delinq.2yrs	pub.rec	not.fully.paid
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	1	0	0


```
[11]: # target value counts
import matplotlib.pyplot as plt
target_value_counts = df['not.fully.paid'].value_counts()
plt.figure(figsize=(6, 6))
plt.bar(target_value_counts.index, target_value_counts.values, color=['pink', 'green'])
plt.title('Not Fully Paid')
plt.xlabel('Status')
plt.ylabel('Count')
plt.xticks(target_value_counts.index, ['Fully Paid', 'Not Fully Paid'])
plt.show()
```



```
[12]: #purpose vs not fully paid
grouped = df.groupby(['purpose', 'not.fully.paid']).size().
    ↪reset_index(name='count')
print(grouped)
fig = px.bar(grouped, x='purpose', y='count', color='not.fully.paid',
    barmode='group', labels={'count': 'Count', 'purpose': 'Purpose',
    ↪'not.fully.paid': 'Not Fully Paid'})
fig.show()
```

	purpose	not.fully.paid	count
0	all_other	0	1944
1	all_other	1	387
2	credit_card	0	1116
3	credit_card	1	146
4	debt_consolidation	0	3354
5	debt_consolidation	1	603

6	educational	0	274
7	educational	1	69
8	home_improvement	0	522
9	home_improvement	1	107
10	major_purchase	0	388
11	major_purchase	1	49
12	small_business	0	447
13	small_business	1	172



```
[13]: df.columns
```

```
[13]: Index(['credit.policy', 'purpose', 'int.rate', 'installment', 'log.annual.inc',
          'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
          'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid'],
          dtype='object')
```

```
[14]: #Relationship between target and features
import seaborn as sns
import matplotlib.pyplot as plt

# List of features
features = ['int.rate', 'installment', 'log.annual.inc', 'dti', 'fico']
# Setting up subplots
fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(14, 12))
fig.suptitle('Relationship between Target and Features')

# Define custom colors
```


```

colors = ['#FF69B4', '#008000']

# Plotting each feature
for i, feature in enumerate(features):
    row = i // 2
    col = i % 2
    sns.barplot(x='not.fully.paid', y=feature, data=df, ax=axes[row, col],
                palette=colors)
    axes[row, col].set_title(feature.capitalize())

# Adjust layout
plt.tight_layout()
plt.show()

```



```

[15]: # Define features and pivot tables
features = ['days.with.cr.line', 'revol.bal', 'revol.util', 'delinq.2yrs', 'inq.
          last.6mths', 'pub.rec']

```



```

grouped = [df.pivot_table(values=feat, index='not.fully.paid', aggfunc='mean')
            ↪for feat in features]


# Setting up subplots
fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(14, 12))
fig.suptitle('Relationship between Target and Features')

# Define custom colors
colors = ['#FF69B4', '#008000']

# Plotting each feature
for i, (feat, group) in enumerate(zip(features, grouped)):
    row = i // 2
    col = i % 2
    sns.barplot(x=group.index, y=feat, data=group.reset_index(), ax=axes[row,
    ↪col], palette=colors)
    axes[row, col].set_title(feat.capitalize())

# Adjust layout
plt.tight_layout()
plt.show()

```



```
[16]: df.columns
```

```
[16]: Index(['credit.policy', 'purpose', 'int.rate', 'installment', 'log.annual.inc',
        'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
        'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid'],
        dtype='object')
```

```
[17]: #DATA ENCODING
```

```
[18]: # Perform one-hot encoding using pandas get_dummies
cat_encoded_df = pd.get_dummies(df['purpose'], drop_first=True)

cat_encoded_df.head() # Displaying the encoded DataFrame
```

```
[18]:   credit_card  debt_consolidation  educational  home_improvement  \
0           0                   1           0           0
1           1                   0           0           0
```

2	0	1	0	0
3	0	1	0	0
4	1	0	0	0

	major_purchase	small_business
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

```
[31]: # Add a suffix to the columns of the encoded DataFrame
cat_encoded_df = cat_encoded_df.add_suffix('_encoded')

# Merge the encoded DataFrame with the original DataFrame
df = pd.concat([df, cat_encoded_df], axis=1)

# Display the updated DataFrame
print(df.head())
```

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	\
0	1	0.1189	829.10	11.350407	19.48	737	
1	1	0.1071	228.22	11.082143	14.29	707	
2	1	0.1357	366.86	10.373491	11.63	682	
3	1	0.1008	162.34	11.350407	8.10	712	
4	1	0.1426	102.92	11.299732	14.97	667	

	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	...	educational	\
0	5639.958333	28854	52.1	0	...	0	
1	2760.000000	33623	76.7	0	...	0	
2	4710.000000	3511	25.6	1	...	0	
3	2699.958333	33667	73.2	1	...	0	
4	4066.000000	4740	39.5	0	...	0	

	home_improvement	major_purchase	small_business	credit_card_encoded	\
0	0	0	0	0	
1	0	0	0	1	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	1	

	debt_consolidation_encoded	educational_encoded	home_improvement_encoded	\
0	1	0	0	
1	0	0	0	
2	1	0	0	
3	1	0	0	
4	0	0	0	

	major_purchase_encoded	small_business_encoded
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 43 columns]

```
[32]: df.columns
```

```
[32]: Index(['credit.policy', 'int.rate', 'installment', 'log.annual.inc', 'dti',
        'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
        'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid',
        'credit_card', 'debt_consolidation', 'educational', 'home_improvement',
        'major_purchase', 'small_business', 'credit_card', 'debt_consolidation',
        'educational', 'home_improvement', 'major_purchase', 'small_business',
        'credit_card', 'debt_consolidation', 'educational', 'home_improvement',
        'major_purchase', 'small_business', 'credit_card', 'debt_consolidation',
        'educational', 'home_improvement', 'major_purchase', 'small_business',
        'credit_card_encoded', 'debt_consolidation_encoded',
        'educational_encoded', 'home_improvement_encoded',
        'major_purchase_encoded', 'small_business_encoded'],
        dtype='object')
```

```
[37]: #Correlation Heatmap.
import plotly.express as px

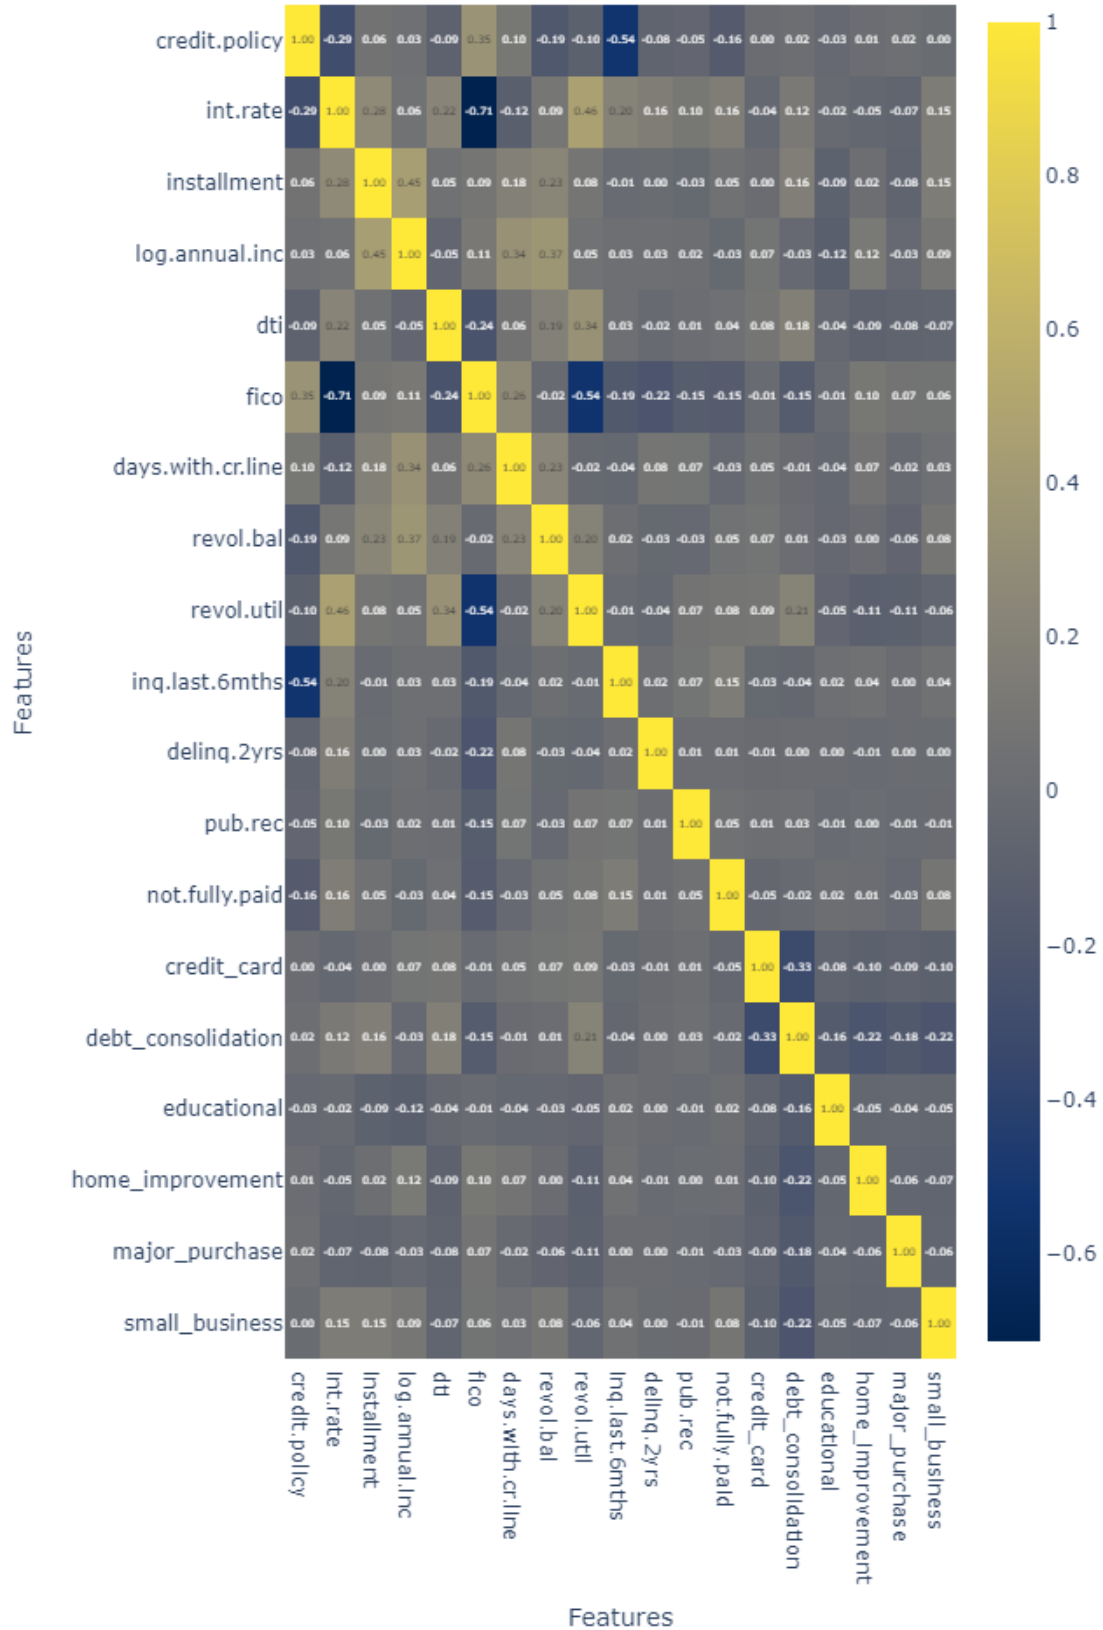
corr = df[['credit.policy', 'int.rate', 'installment', 'log.annual.inc',
        'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
        'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid',
        'credit_card', 'debt_consolidation',
        'educational', 'home_improvement',
        'major_purchase', 'small_business']].corr()

fig = px.imshow(corr, text_auto=".2f", color_continuous_scale='Cividis',
        aspect="auto")

fig.update_layout(title='Correlation Heatmap',
        xaxis_title='Features',
        yaxis_title='Features')
fig.update_layout(height=1000)

fig.show()
```

Correlation Heatmap



```
[ ]: #DATA SPLITTING
```

```
[38]: df.columns
```

```
[38]: Index(['credit.policy', 'int.rate', 'installment', 'log.annual.inc', 'dti',  
        'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',  
        'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid',  
        'credit_card', 'debt_consolidation', 'educational', 'home_improvement',  
        'major_purchase', 'small_business', 'credit_card', 'debt_consolidation',  
        'educational', 'home_improvement', 'major_purchase', 'small_business',  
        'credit_card', 'debt_consolidation', 'educational', 'home_improvement',  
        'major_purchase', 'small_business', 'credit_card', 'debt_consolidation',  
        'educational', 'home_improvement', 'major_purchase', 'small_business',  
        'credit_card_encoded', 'debt_consolidation_encoded',  
        'educational_encoded', 'home_improvement_encoded',  
        'major_purchase_encoded', 'small_business_encoded'],  
        dtype='object')
```

```
[42]: #defining features and target variable  
x = df[['credit.policy', 'int.rate', 'installment', 'log.annual.inc', 'dti',  
        'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',  
        'inq.last.6mths', 'delinq.2yrs', 'pub.rec',  
        'credit_card', 'debt_consolidation',  
        'educational', 'home_improvement',  
        'major_purchase', 'small_business']]  
y = df['not.fully.paid']  
print(len(x))  
print(len(y))
```

```
9578
```

```
9578
```

```
[43]: #splitting the data into training and testing.  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25)  
print(x_train.shape)  
print(x_test.shape)
```

```
(7183, 36)
```

```
(2395, 36)
```

```
[45]: #scaling the data sets  
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
x_train_scaled = scaler.fit_transform(x_train)  
x_test_scaled = scaler.transform(x_test)
```

```
[ ]: #CUSTOM NEURAL NETWORK
```

```
[46]: #custom model with hidden layers
import tensorflow as tf
class CustomNetwork(tf.keras.Model):
    def __init__(self):
        super(CustomNetwork, self).__init__()
        self.dense1 = tf.keras.layers.Dense(100, activation='sigmoid')
        self.dropout1 = tf.keras.layers.Dropout(0.2)
        self.dense2 = tf.keras.layers.Dense(50, activation='sigmoid')
        self.dropout2 = tf.keras.layers.Dropout(0.2)
        self.dense3 = tf.keras.layers.Dense(20, activation='sigmoid')
        self.dropout3 = tf.keras.layers.Dropout(0.2)
        self.dense4 = tf.keras.layers.Dense(1, activation='sigmoid')

    def call(self, inputs):
        x = self.dense1(inputs)
        x = self.dropout1(x)
        x = self.dense2(x)
        x = self.dropout2(x)
        x = self.dense3(x)
        x = self.dropout3(x)
        x = self.dense4(x)
        return x

model = CustomNetwork()
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.01),
              loss=tf.keras.losses.BinaryCrossentropy(), #binary cross entropy
              ↪because the classification is binary 1,0
              metrics=['accuracy'])

model.fit(x_train_scaled, y_train, epochs=10)
```

Epoch 1/10

225/225 [=====] - 1s 1ms/step - loss: 0.4478 - accuracy: 0.8338

Epoch 2/10

225/225 [=====] - 0s 1ms/step - loss: 0.4302 - accuracy: 0.8359

Epoch 3/10

225/225 [=====] - 0s 1ms/step - loss: 0.4263 - accuracy: 0.8359

Epoch 4/10

225/225 [=====] - 0s 1ms/step - loss: 0.4251 - accuracy: 0.8359

Epoch 5/10

```

225/225 [=====] - 0s 988us/step - loss: 0.4205 -
accuracy: 0.8359
Epoch 6/10
225/225 [=====] - 0s 975us/step - loss: 0.4214 -
accuracy: 0.8356
Epoch 7/10
225/225 [=====] - 0s 973us/step - loss: 0.4184 -
accuracy: 0.8359
Epoch 8/10
225/225 [=====] - 0s 982us/step - loss: 0.4184 -
accuracy: 0.8354
Epoch 9/10
225/225 [=====] - 0s 984us/step - loss: 0.4175 -
accuracy: 0.8346
Epoch 10/10
225/225 [=====] - 0s 980us/step - loss: 0.4170 -
accuracy: 0.8368

```

```
[46]: <keras.src.callbacks.History at 0x7f2fcdeaca90>
```

```
[47]: #model predictions
y_pred = model.predict(x_test_scaled)
y_pred
```

```
75/75 [=====] - 0s 624us/step
```

```
[47]: array([[0.4469795 ],
            [0.2629432 ],
            [0.13791387],
            ...,
            [0.25922832],
            [0.1283481 ],
            [0.13250448]], dtype=float32)
```

```
[49]: df = pd.DataFrame({'Prediction': y_pred.flatten(), 'Label': y_test})
print(df.head(5))
```

	Prediction	Label
8465	0.446979	1
3270	0.262943	0
7356	0.137914	0
1380	0.227290	0
2607	0.130376	0

```
[50]: #model evaluation
accuracy = accuracy_score(y_test, y_pred.round())
print("Accuracy:", accuracy)
```



Accuracy: 0.8496868475991649

```
[52]: #Confusion Matrix.
import plotly.express as px
# Assuming y_test and y_pred are defined
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred.round())

fig = px.imshow(cm, text_auto=True, color_continuous_scale='Pinky1',
                labels=dict(x='Predicted', y='True', color='Count'),
                title='Confusion matrix')

fig.show()
```

Confusion matrix



```
[53]: model2 = tf.keras.Sequential([
    tf.keras.layers.Dense(100, activation='sigmoid'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(50, activation='sigmoid'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(20, activation='sigmoid'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model2.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.01),
               loss=tf.keras.losses.BinaryCrossentropy(),
               metrics=['accuracy'])
```

```
model2.fit(x_train_scaled, y_train, epochs=10)
```

```
Epoch 1/10
225/225 [=====] - 1s 1ms/step - loss: 0.4447 -
accuracy: 0.8332
Epoch 2/10
225/225 [=====] - 0s 996us/step - loss: 0.4318 -
accuracy: 0.8359
Epoch 3/10
225/225 [=====] - 0s 995us/step - loss: 0.4275 -
accuracy: 0.8359
Epoch 4/10
225/225 [=====] - 0s 971us/step - loss: 0.4238 -
accuracy: 0.8359
Epoch 5/10
225/225 [=====] - 0s 984us/step - loss: 0.4193 -
accuracy: 0.8353
Epoch 6/10
225/225 [=====] - 0s 981us/step - loss: 0.4214 -
accuracy: 0.8359
Epoch 7/10
225/225 [=====] - 0s 988us/step - loss: 0.4195 -
accuracy: 0.8360
Epoch 8/10
225/225 [=====] - 0s 989us/step - loss: 0.4152 -
accuracy: 0.8347
Epoch 9/10
225/225 [=====] - 0s 990us/step - loss: 0.4148 -
accuracy: 0.8345
Epoch 10/10
225/225 [=====] - 0s 978us/step - loss: 0.4166 -
accuracy: 0.8352
```

```
[53]: <keras.src.callbacks.History at 0x7f2f2c1eb100>
```

```
[54]: y_pred2 = model2.predict(x_test_scaled)
      y_pred2
```

```
75/75 [=====] - 0s 640us/step
```

```
[54]: array([[0.40044606],
          [0.26314956],
          [0.10542049],
          ...,
          [0.25139442],
          [0.13323942],
```

```
[0.1449118 ]], dtype=float32)
```

```
[55]: #Printing accuracy.  
accuracy = accuracy_score(y_test, y_pred2.round())  
print("Accuracy:", accuracy)
```

```
Accuracy: 0.8521920668058455
```

```
[ ]:
```